



Computer Graphics:

Introduction to Shadertoy

By: Nick | ACM at UCSD



What is Computer Graphics?

Sorry to those who have taken 167...

Here's the generic Wikipedia definition:

“Computer graphics is the discipline of generating images with the aid of computers”

Computer graphics is used in a variety of applications, mainly:

- Film (either for animations or for visual effects)
- Video Games



UNREAL
ENGINE



Animation



Visual Effects

Shaders




If you have experience drawing things with computers, you may know that things are drawn *sequentially*

For example, first draw a circle here, a line here, etc until you finish drawing your object

Shaders are also a set of instructions, except the instructions are executed *all at once* (but not necessarily for every pixel).

One dot at a time



All at once

ShaderToy



ShaderToy sets up the shader so that it runs once for each pixel

Acts like a function that receives the pixel position and a few other inputs and returns the pixel color as an output

For those who have taken 167, it basically rasterizes a full-screen quad and the shader you write is the fragment shader. It also uses WebGL (as opposed to OpenGL, though it is a similar standard) to allow it to run in a browser

GLSL



OpenGL shaders are written in a language called GLSL
*open**GL** Shading **L**anguage*

Has many built-in features as opposed to regular programming languages, mainly support for basic Linear Algebra

Oh yes we will be doing some math

PLEASE DON'T FREAK OUT JUST YET PLEASE IT'S JUST BASIC
LINEAR ALGEBRA



Image

Shader Inputs

```
1  const float PI = 3.14159265;
2
3  void mainImage( out vec4 fragColor, in vec2 fragCoord )
4  {
5      // Normalized pixel coordinates (from 0 to 1)
6      vec2 uv = fragCoord/iResolution.xy;
7
8      float y = 0.5 + 0.5 * sin(uv.x * 2.0 * PI);
9
10     // calculate color here
11
12     // Output to screen
13     fragColor = vec4(col,1.0);
14 }
```




My First ShaderToy



Now calculate the color!

Simple sin graph (answer):

<https://www.shadertoy.com/view/wdlcDl>



REC



0

Created by **nickwn** in 2020-04-03



Image

Shader Inputs

```
1  const float PI = 3.14159265;
2
3  void mainImage( out vec4 fragColor, in vec2 fragCoord )
4  {
5      // Normalized pixel coordinates (from 0 to 1)
6      vec2 uv = fragCoord/iResolution.xy;
7
8      float y = 0.5 + 0.5 * sin(uv.x * 2.0 * PI);
9
10     vec3 col = vec3(0.0);
11     if(uv.y < y)
12     {
13         col = vec3(1.0);
14     }
15
16     // Output to screen
17     fragColor = vec4(col,1.0);
18 }
```

Nice, but one problem



Some aliasing issues

There's a branch (an if statement), which isn't that bad but it gives us an opportunity to learn how the GPU works

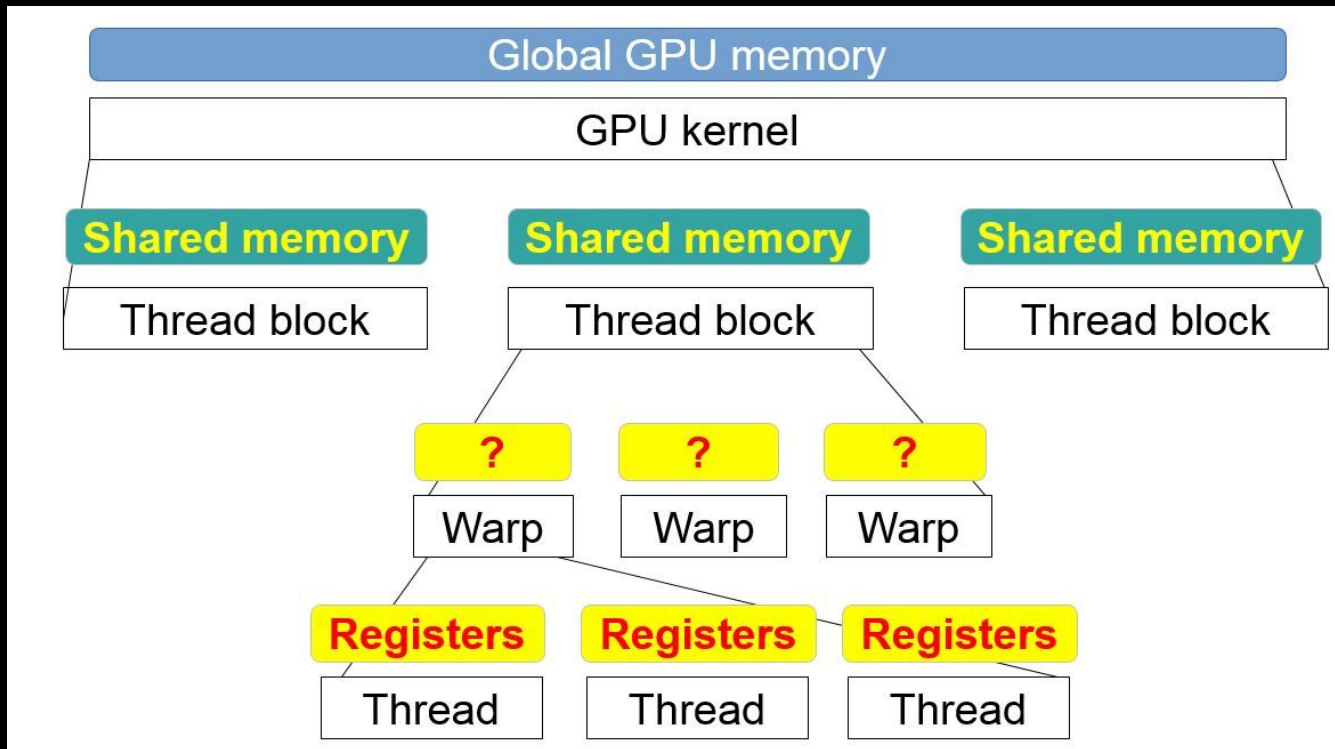


Why is this a problem?

To better understand, we need to look at the basic architecture of a GPU (terminology may differ):

- GPUs are filled with hundreds and hundreds of processing “threads” which run your code in parallel
- These processing threads are separated into small(ish) groups, called *warps*, which are given work by schedulers

A Fun Diagram





Branching in Warps

The problem is that a warps are issued a single instruction at a time

This means that all threads in a warp have to run the same instruction together

This is inherently against how branches work; ideally threads should be executing different instructions at the same time based on a condition

How does the GPU solve this?




More Fun Diagrams

Here's an example of
your typical CPU
control flow:

If condition is true, the
green path is executed.

If condition is false, the
red path is executed






More Fun Diagrams

Here's how it looks on the GPU (where each arrow represents a thread):

The GPU basically creates a mask for all the threads based on the condition


The ALUs execute every instruction, but their output is masked out if they are unneeded



More Fun Diagrams



Execution within warps also removes one branch if it can, so if only one side of the branch needs to be executed for all threads, then the other side will not be executed



Divergence



The idea that not all the work threads do is productive is important to optimization

If it is very likely for a thread to not be doing productive work, your program has *high divergence*

Applies to branching, but also applies to things like loops. An entire warp could be blocked if it needs to wait for a single thread to complete its loop.



How do we remove the if?

We can remove the `if` using the `smoothstep` function in glsl

`smoothstep` is basically a fancy interpolation function that can be used to smoothly blend between two values, in our case black and white:


Key thing to note is that it doesn't actually do the interpolation; it just returns the t-value for the interpolation equation (to interpolate between a and b):

$$interp = a * (t-1) + b * t$$

smoothstep



<https://thebookofshaders.com/glossary/?search=smoothstep>



My First ShaderToy (take 2)



Modify your previous code to remove the if statement!

Better sin graph (answer):

<https://www.shadertoy.com/view/tslcDl>



Created by **nickwn** in 2020-04-03



Image

Shader Inputs

```
1  const float PI = 3.14159265;
2
3  void mainImage( out vec4 fragColor, in vec2 fragCoord )
4  {
5      // Normalized pixel coordinates (from 0 to 1)
6      vec2 uv = fragCoord/iResolution.xy;
7
8      float y = 0.5 + 0.5 * sin(uv.x * 2.0 * PI);
9
10     float percent = smoothstep(y + 0.02, y, uv.y);
11
12     vec3 col = vec3(percent);
13
14     // Output to screen
15     fragColor = vec4(col,1.0);
16 }
```

Raymarching



Bonus section

aboutta go 0-100 real quick



Signed Distance Functions

A function that, when passed the coordinates of a point in space, returns the shortest distance between that point and some surface

The sign of the return value indicates whether the point is inside that surface or outside


For example, the SDF for a sphere looks like this:

$$f(x, y, z) = \sqrt{x^2 + y^2 + z^2} - r$$

Raymarching



We can “step” along a camera view ray at different distances, evaluating the SDF at each step to see if we are close to a surface and to determine the distance to the next step.



Resources



<http://iquilezles.org/www/index.htm>

- SDFs:
<http://iquilezles.org/www/articles/distfunctions/distfunctions.htm>
- <http://iquilezles.org/www/articles/distfunctions2d/distfunctions2d.htm>
- Examples:
<http://iquilezles.org/www/articles/raymarchingdf/raymarchingdf.htm>

Final Project



Color in the surfboard/add a beach ball:

<https://www.shadertoy.com/view/wdScRz>

Final Project



For the Future



Raymarching as it is done in shadertoy is kinda niche

As a general concept however, it is used extensively in volume rendering and a little for global illumination

Download Blender for the workshop next week, where we will be creating a new material using Blender's material graph and OSL that can approximate sub-surface scattering effects

Sub-org???: acmurl.com/graphics-interest



Slide References

UPenn SIGGRAPH's Shadertoy presentation:

<https://drive.google.com/file/d/1XX4OMwRxrMcgU5NTBPz2qjRKOhRAcPhp/view?usp=sharing>

Mythbusters: <https://www.youtube.com/watch?v=ZrJeYExpUyQ>

The Book of Shaders: <https://thebookofshaders.com/>

Old Nvidia presentation:

<https://www.nvidia.com/docs/IO/116711/sc11-perf-optimization.pdf>

New Nvidia Post: <https://devblogs.nvidia.com/register-cache-warp-cuda/>